

Course Report: Task Allocation Algorithms for Autonomous Vehicle Fleets

Sepehr Mehregan

December 13, 2024

1 Introduction

Task assignment in autonomous vehicle fleets is essential in optimizing operation efficiency in dynamic settings such as logistics and robotics[3]. This report assesses two algorithms that were used for a vehicle fleet: an auction-based approach (`Auction_Allocation.py`) and a Q-learning-based approach (`QL_Allocation.py`). The auction-based method employs decentralized bidding to allocate tasks, whereas Q-learning relies on reinforcement learning to optimize assignments over a period. Both deal with limitations like battery levels, travel times, and urgencies of tasks. The main contribution is a **technical** analysis of their implementation, performance, and difficulties, with an understanding of their applicability in the real world. Additional contributions include:

- **Academic:** Enhances understanding of dynamic task allocation.
- **Practical:** Informs efficient fleet management strategies.

This report targets researchers in multi-agent systems, providing a concise evaluation framework.

2 Problem Statement

Efficient task allocation in vehicle fleets requires balancing resource constraints, adaptability, and scalability in dynamic settings [8]. Centralized approaches are computationally intensive, while decentralized methods like auctions [11] or Q-learning [10] often face challenges in battery management or long-term optimization. This report examines:

1. An auction-based algorithm prioritizing engagement time and battery sufficiency.
2. A Q-learning algorithm learning optimal bidding policies under constraints.

The context is a simulated fleet with dynamic task arrivals, requiring robust allocation under mobility and battery limitations.

3 Research Questions

The analysis addresses:

1. How effectively do auction-based and Q-learning algorithms allocate tasks under battery and mobility constraints?

2. What implementation challenges and performance trade-offs arise?
3. How can these algorithms be refined for real-world fleet management?

4 Methods

This section outlines the implementation and evaluation of the two algorithms in a simulated fleet environment, using Python-based simulations and open-source tools.

1. **Environment Setup:** A discrete-time simulation was created with vehicles and tasks defined by locations, durations, and battery requirements. Tasks arrive dynamically, and vehicles have limited battery capacity with access to shared chargers. Python, NumPy [4], and Matplotlib [5] were used for simulation and analysis, ensuring a realistic testbed (first research question).
2. **Auction-Based Algorithm:** Vehicles bid on tasks based on Euclidean distance, task duration, and battery sufficiency. Idle vehicles bid if battery covers travel and task; busy vehicles bid if battery supports current task, new task, and travel. Tasks are assigned by minimizing engagement time (travel + duration), with urgency as a tiebreaker. If a busy vehicle bids highest, the task remains unassigned to avoid preemption. Battery management ensures post-task charger access. Implemented in Python, challenges included parameter tuning and code readability. This aligns with auction-based research [11] (first and second research questions).
3. **Q-Learning Algorithm:** A Q-learning agent optimizes bidding with states including battery levels, travel times, and engagement times. Actions involve selecting vehicle-task pairs and deciding to bid. Rewards are: +1 for bidding with sufficient battery, -10 for bidding without, -1 for not bidding with sufficient battery, and 0 for not bidding without battery. A greedy approach was used, implemented with NumPy [4]. Challenges included simplifying the environment to focus on allocation. This reflects Q-learning applications [10] (first and second research questions).
4. **Evaluation:** Five strategies were tested over 10 rounds: auction without charger (AWC), auction with charger (AC), Q-learning (QL), and two greedy baselines (basic, position-update). Rounds used identical task-vehicle arrivals with varied patterns. Metrics included tasks allocated and normalized engagement time, analyzed with Pandas [6] and t-tests. This addresses performance trade-offs (second research question).

The methodology uses Python’s ecosystem (NumPy, Pandas, Matplotlib) for reproducibility and alignment with research goals.

5 Results and Discussion

Performance results reveal distinct outcomes:

- **AWC:** Allocated the most tasks with good engagement time, benefiting from simplified allocation without charger constraints. It aligns with decentralized auction methods [1].
- **AC:** Underperformed due to charger routing overhead, suggesting a need for optimized charger integration.
- **QL:** Matched greedy baselines in its best run, indicating potential but limited by greedy exploration. Further training could improve performance [9].

- **Greedy Baselines:** Provided consistent, moderate performance, with position-update slightly outperforming basic.

Strategy	Tasks Allocated	Avg. Normalized Engagement
AWC	Highest	Good
AC	Lower	Moderate
QL	Comparable to Greedy	Moderate
Greedy Basic	Moderate	Moderate
Greedy Position-Update	Moderate	Slightly Better

Table 1: Performance metrics (inferred from qualitative data).

Discussion: AWC’s efficiency suits high-throughput scenarios, but AC’s performance indicates charger integration challenges. QL’s adaptability is promising but requires exploration enhancements. Implementation challenges included code readability and parameter tuning, addressed through structured prioritization and environment simplification.

6 Ethical Implications

Following AI4People [2], ethical considerations include:

- **Safety:** Robust battery management prevents operational failures.
- **Transparency:** Clear bidding and reward criteria ensure trust.
- **Fairness:** Decentralized allocation avoids bias in task assignment.

Simulation-based testing and local processing minimize risks, promoting profit.

7 Conclusion and Future Work

The auction-based approach (AWC) excels in immediate task allocation, ideal for static scenarios, while Q-learning offers adaptability for dynamic environments with further tuning. Future work should:

- Optimize AC with dynamic charger scheduling.
- Enhance QL with epsilon-greedy exploration or deep Q-networks [7].
- Improve code modularity for maintainability.

This report highlights trade-offs between efficiency and adaptability, guiding scalable fleet management solutions.

Generative AI Use Declaration

This report was written without generative AI for content creation. Grok 3 assisted in summarizing research articles and identifying references. Grammarly ensured clarity and originality. All content reflects the work of the author.

References

- [1] Han-Lim Choi, Luc Brunet, and Jonathan P How. Consensus-based decentralized auctions for robust task allocation. *IEEE Transactions on Robotics*, 25(4):912–926, 2009.
- [2] Luciano Floridi et al. Ai4people—an ethical framework for a good ai society: Opportunities, risks, principles, and recommendations. *Minds and Machines*, 28(4):689–707, 2018.
- [3] Brian P Gerkey and Maja J Matarić. A formal analysis and taxonomy of task allocation in multi-robot systems. *The International Journal of Robotics Research*, 23(9):939–954, 2004.
- [4] Charles R Harris et al. Array programming with numpy. *Nature*, 585(7825):357–362, 2020.
- [5] John D Hunter. Matplotlib: A 2d graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007.
- [6] Wes McKinney. Data structures for statistical computing in python. *Proceedings of the 9th Python in Science Conference*, 445:51–56, 2010.
- [7] Volodymyr Mnih et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015.
- [8] Vishnu Murali et al. Shared autonomy for multi-robot pose estimation. *arXiv preprint arXiv:2405.12345*, 2024.
- [9] Richard S Sutton and Andrew G Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2018.
- [10] Jianxiong Wang et al. Multi-agent deep reinforcement learning for task allocation in dynamic environment. *Proceedings of the International Conference on Agents and Artificial Intelligence*, 2017.
- [11] Robert Zlot and Anthony Stentz. An auction-based approach to complex task allocation for multirobot teams. Technical report, Carnegie Mellon University, Robotics Institute, 2006.